

# ENVIRONMENTAL SOUND CLASSIFICATION WITH SEMI-SUPERVISED LEARNING

*Rasmus Arpe Fogh Jensen (s134843), Thomas Pethick (s144448)*

Technical University of Denmark

## ABSTRACT

In this paper, we explore a semi-supervised learning approach for environmental sound classification. A convolutional autoencoder is used for pre-training the weights in the network. Two different methods for the inversions of the max-pooling layers are examined in the decoder; up-sampling and unpooling. The semi-supervised approach is benchmarked against a supervised approach with similar architecture on a public available dataset. The results show that semi-supervised learning yields slightly better performance utilizing unlabeled data.

**Index Terms**— Environmental sound, classification, semi-supervised learning, convolutional neural networks

## 1. INTRODUCTION

The application of deep learning methods in the context of speech and audioprocessing has recently gained a lot of attention. The models often rely on a log-scaled mel-spectrogram representation of audio samples. Google has however shown that it is possible to obtain state of the art performance for some audioprocessing tasks using convolutions directly on the raw audio feed with their proposed Wavenet model [1].

Research in classification of environmental sounds (everyday audio events that do not consist of music or speech) has been performed with results showing that deep learning methods are able to outperform previously preferred methods such as unsupervised feature learning for a variety of datasets [2, 3]. The state-of-the-art models perform convolutions on the mel-spectrogram representation of the sound data with a fully connected layer at the top and also utilize feature engineering by performing various data-augmentations [3].

A common problem with deep learning for sound classification is the lack of labeled data. Google has recently introduced the Audioset [4] dataset in order to combat the problem and the dataset can be exploited in transfer-learning methods. When the domain of the sounds is substantially different from the Audioset, transfer learning might not yield the best results. Companies often have unlabeled data available for their particular domain of interest. An interesting field of

research is therefore how we can utilize the unlabeled data for pre-training in order to achieve better classification results.

In this article, we explore the performance of a semi-supervised learning approach on the environmental sound classification task. We use convolutional autoencoders as a way of pre-training the weights in the network to obtain better discrimination. Two different structures for the decoding part of the autoencoder are explored; upsampling and unpooling. We benchmark our semi-supervised method against a convolutional architecture with dense layers on top. The results are obtained using public available datasets.

The article is organized as following. Section 2 describes our proposed approach to the problem. Section 3 describes the datasets used and the pre-processing required. Section 4 explains how the experiments were conducted and the hyperparameter optimization. Section 5 shows the results of the approach. At last, we have a discussion of the method and results in section 6.

### 1.1. Related Work

Environmental Sound classification has been done by many - Augmenting data - Piczak

Semi-supervised approaches - Denoising autoencoder - Google wavenet - Unet?

## 2. METHOD

### 2.1. Semi-supervised learning

Semi-supervised learning is a class of supervised learning where training makes use of unlabeled data in addition to the labeled data. Useful features can be learned using the unlabeled data that can improve classification.

More precisely training is done on  $n$  unlabeled examples  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$  and  $m$  examples  $\mathbf{x}^{(n+1)}, \dots, \mathbf{x}^{(n+m)}$  with corresponding labels  $\mathbf{y}^{(n+1)}, \dots, \mathbf{y}^{(n+m)}$ . The goal is to achieve a higher classification accuracy using both datasets than what would be possible by discarding the unlabeled data and doing supervised learning.

---

Thanks to Corti for supervising the project

## 2.2. Autoencoder

An autoencoder is suitable for semi-supervised learning. It consists of an encoder and a decoder where the encoder maps the input to a hidden representation and the decoder then reconstructs the input based on that hidden representation [5]. To be precise, the autoencoder takes an input  $\mathbf{x} \in \mathbb{R}^d$  encodes it to  $\mathbf{z} \in \mathbb{R}^{d'}$  using a deterministic function  $\mathbf{z} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$ . The input is then reconstructed using a reverse mapping  $\tilde{\mathbf{x}} = \sigma(\mathbf{W}'\mathbf{z} + \mathbf{b}')$ .

The weights are trained by minimizing the difference between  $\mathbf{x}$  and  $\tilde{\mathbf{x}}$  called the reconstruction error[6],

$$E_{rec} = \frac{1}{2n} \sum_{i=1}^n \|\tilde{\mathbf{x}}^{(i)} - \mathbf{x}^{(i)}\|^2.$$

To avoid that the autoencoder simply learn the identity function some regularization is needed. These methods involve introducing a regularizing term in the cost function that imposes sparsity. We will however enforce sparsity through the network architecture as suggested by [7] and covered in section 2.3.

The autoencoder will be used in the semi-supervised learning approach by providing weight initialization for part of the final classification network. First the autoencoder is trained on the unlabeled data. The encoder is then extended with a classifier and trained on the labeled data. This final step is called *fine-tuning* and adjusts all weights in the network. We will refer to the classifier initialized with weights from the autoencoder as being *pre-trained*.

## 2.3. Convolutional Autoencoder (CAE)

Instead of a fully connected autoencoder, the encoder can consist of convolutions as proposed by [7]. The main observation is that a max-pooling layer in the hierarchical network can be used to impose sparsity so the encoder does not learn a trivial identity function even if the hidden representation is of higher dimensionality than the input. By having the max-pooling layer, the need for regularizers like L1 and L2 are avoided.

The authors also experimented with adding noise to the input as a regularizing effect (but of course still evaluating the cost function on the original input). However, this did not seem to lead to much more biological plausible filters compared to max-pooling, so we have avoided this additional complexity.

Originally the authors did the unsupervised pre-training of the stacked convolutional autoencoder in a greedy, layer-wise fashion. Subsequently, good results has been achieved on networks as deep as the ones we consider by training the whole network at once [8]. Consequently, we have decided to train all weights simultaneously.

## 2.4. Unpooling

The question of how to construct the decoder arises when the encoder is a combination of convolutions and max-pooling layers.

To create a function mapping from  $\mathbf{z}$  to  $\mathbf{x}$ , we reverse all layers. A transposed convolution is commonly used in the decoder as originally proposed in [9]. We refrain from using the term deconvolution as it is not the inverse of convolution. For a thorough explanation of the differences we refer to [10] [11].

It is not however obvious how max-pooling should be reversed. In the encoder, downsampling occurred whereas the decoder has to upsample. Max-pooling is a many-to-one relationship and the location in the input of the picked output value is not necessarily known when doing upsampling. The naive approach uses the input by repeating each value to neighboring fields in the output which we will refer to as *upsampling*.

Zhao et al. [12] proposes a different solution that records the position used in the max-pooling layer. Using this position the corresponding layer in the decoder then only copies the input value to the output on that location. All other output values are set to a value of 0. The authors argues that this approach, which we will refer to as *unpooling*, provides a regularization effect and thus leads to better generalization and reconstructions.

Since an implementation neither existed in tensorflow, theano or Keras at the time of writing we have made an implementation of a unpooling layer in Keras available at <https://github.com/...>. However, note that this layer is currently only supported by the theano backend since it requires taking the gradient of the gradient which is not available in Tensorflow.

## 3. DATA

All of the results are obtained using the datasets ESC-50 and ESC-US [13].

ESC-50 is a collection of 2000 labeled environmental sounds manually extracted from public field recordings gathered from Freesound.org. The dataset includes sounds such as animal sounds, water sounds and urban sounds. Baseline classifiers using *k-nearest neighbour* (K-NN) and *random Forest* (RF) performs with 32% and 45% accuracy, respectively. However, [2] obtained an accuracy of 62% using a convolutional neural network approach.

ESC-US is a collection of 250.000 unlabeled sounds also extracted from public field recordings of Freesound.org. They are however not manually annotated. ESC-US is suited for unsupervised pre-training.

All of the sounds from ESC-50 and ESC-US datasets are of 5 second length sampled at 44.1kHz.

### 3.1. Pre-processing

The sounds were normalized and processed into the frequency domain of log-scaled mel-spectrograms using the Librosa[14] python library. The mel-spectrograms were extracted using 28 mel-bands, window size of 12288 and hop-length of 6144. This gives in an input dimension of  $\mathbf{x}^{(n)} \in \mathbb{R}^{28,36}$ . The mel-spectrogram transformation involves a small mel-bandwidth and a large window size compared to results from [2, 3], hence some features important for classification might not be present. The choice of these parameters is based on faster training and iteration since the time scope for the project was limited. However, the architecture and the pre-processing is easily scalable for future experimentation.

The Pescador[15] library is used to handle streams of data fed to the neural networks.

## 4. EXPERIMENTATION

We adopt the notation used in [12] to describe the considered architectures. The notation  $(64)5c-2p-50fc$  should be read as a 5x5 convolutional kernel with 64 filters, followed by a 2x2 max-pooling layer, and finally ending with a fully connected layer with 50 hidden units. When used to specify an autoencoder, only the encoder is described while the decoder is implied by that specification. Unpooling is assumed in the decoder when nothing else is explicitly stated.

### 4.1. Network setup

The Relu activation function is used throughout the network except in the last layer of the decoder where a linear activation is needed for the mean squared error function. The network is trained using a batch size of 32.

- Learning rate, Activation also softmax, Batches drawn

### 4.2. Early Stopping

In addition the the regularization effect of max-pooling we employ early stopping to avoid over-fitting. This is a method that stops the training loop when the validation loss no longer decreases. A patience of 5 is used to allow for fluctuations especially in the beginning before it starts converging.

### 4.3. Reconstructions

### 4.4. Architecture

The architecture is inspired by the VGG-16 based autoencoder used for semantic segmentation which leverages unpooling [16]. However, it adopts the architecture used for MNIST in original paper on unpooling by Zhao et al. [7] since the input size is in the same order of magnitude. The small squared convolutional kernels used in this architecture are performant even for mel-spectrograms [17] [18].

The base architecture used for the autoencoder was  $(f)5c-2p-(f)3c-2p-(f)3c-2p$  with batch size 32 using unpooling. Exhaustive grid search was employed for batch sizes 16, 32, 64 and number of filters  $f \in \{16, 32, 64\}$ . The biggest filter size 64 yielded better results in all cases. This is expected so long we regularize properly since the model is more expressive. Even bigger filters where not consider because it would extend training time.

The remaining parameters was chosen heuristically due to the extensive training time required. The Adam optimizer was used with a learning rate of 0.001 after observing that 0.01 did not converge after several epochs. We experimented with a smaller network of only 2 convolutional layers but this yielded a dramatic decrease in performance so stopped pursuing that direction.

Upsampling was quickly discarded. A pretrained CNN network using CAE with upsampling achieve a classification accuracy of 0.018 after 7 epochs of pretraining and 29 epochs of classification training. Without pre-training the same network achieved an accuracy of 0.41 in 25 epochs. Considering this is approximately the prediction accuracy of randomly guessing this seems suspicious. It might be stuck in some local optimum where the learning rate is not big enough to adjust the weights sufficiently fast for it to learn over the 30 epochs.

Overfitting did seem to be an issue with classification accuracy reaching 1 for the best performing models. Batch normalization was tested as regularization to avoid overfitting. However, this did not have positive results. Possible explanations for this will be given in section 5.


## 5. RESULTS

The final performance was improved as is evident from table 1. This aligns with the findings that unsupervised pre-training helps by supporting better generalization achieved through guiding the learning towards basins of minima [19].


Plotting the accuracy during training showed that for all tried combination of parameters the training accuracy was converging with the validation accuracy always lower but steadily increasing. We noticed that most of the pre-trained classifiers started with a validation accuracy around 0.1. This initial boosted accuracy is expected by having part of the networks weights cleverly initialized.

## 6. CONCLUSION

The test data set might be too similar to training set since the accuracy is quite high even without regularization methods like dropout.



**Fig. 1.** Display of reconstruction on test set using (64) 5c-2p- (64) 3c-2p- (64) 3c-2p-50fc. From top to bottom: 1) ground truth 2) reconstruction using unpooling 3) reconstruction using upsampling.



**Fig. 2.** Final CNN classifier Architecture

Architecture		Accuracy	
b	f	Supervised	Pre-training with CAE
16	64	0.414	0.502
32	64	0.492	<b>0.507</b>
64	64	0.474	-

**Table 1.** Accuracy results for architecture (f) 5c-2p- (f) 3c-2p- (f) 3c-2p where f is the number of filters and b is the batch size.

Network	Accuracy
Baseline Random Forest	44%
Piczak CNN	65.5%
Our CAE pre-trained CNN	50.7%

**Table 2.** Accuracy for pre-trained (64) 5c-2p- (64) 3c-2p- (64) 3c-2p-128fc-50fc network compared against baseline and CNN results obtained by Piczak [2].

## 6.1. Future work

We expect the small input size to be the main course for the lack of performance. Usually a much higher band size is used as well as smaller window size leading to more frames. It is highly likely that the smaller spectrograms will lack crucial features from the original waveform input necessary to distinguish between classes. However, the proposed architecture has the potential of scaling to bigger input sizes. This would be the natural next step when considering how the accuracy could be improved.

Unpooling is not necessarily better than upsampling for learning representation suitable for classification: <https://arxiv.org/pdf/1701>

Joined approach for Semi-Supervised learning with Autoencoders (provides good intro to autoencoders) <http://ieeexplore.ieee.org/> "deep" CNN for spectrograms: <http://www.inf.ufpr.br/lesoliveira/download>

## 7. REFERENCES

- [1] Aäron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves,




Fig. 3. Final CAE Architecture

- Nal Kalchbrenner, Andrew W. Senior, and Koray Kavukcuoglu, “Wavenet: A generative model for raw audio,” *CoRR*, vol. abs/1609.03499, 2016.
- [2] Karol J. Piczak, “Esc: Dataset for environmental sound classification,” in *Proceedings of the 23rd ACM International Conference on Multimedia*, New York, NY, USA, 2015, MM ’15, pp. 1015–1018, ACM.
- [3] Justin Salamon and Juan Pablo Bello, “Deep convolutional neural networks and data augmentation for environmental sound classification,” *CoRR*, vol. abs/1608.04363, 2016.
- [4] Jort F. Gemmeke, Daniel P. W. Ellis, Dylan Freedman, Aren Jansen, Wade Lawrence, R. Channing Moore, Manoj Plakal, and Marvin Ritter, “Audio set: An ontology and human-labeled dataset for audio events,” in *Proc. IEEE ICASSP 2017*, New Orleans, LA, 2017, Retrieved October 25, 2017 from <https://static.googleusercontent.com/media/research.google.com/da//pubs/archive/45857.pdf>.
- [5] Wei Luo, Jun Li, Jian Yang, Wei Xu, and Jian Zhang, “Convolutional sparse autoencoders for image classification,” *IEEE transactions on neural networks and learning systems*, 2017.
- [6] Wu Haiyan, Yang Haomin, Li Xueming, and Ren Haijun, “Semi-supervised autoencoder: A joint approach of representation and classification,” in *Computational Intelligence and Communication Networks (CICN), 2015 International Conference on*. IEEE, 2015, pp. 1424–1430.
- [7] Jonathan Masci, Ueli Meier, Dan Cireşan, and Jürgen Schmidhuber, “Stacked convolutional auto-encoders for hierarchical feature extraction,” *Artificial Neural Networks and Machine Learning–ICANN 2011*, pp. 52–59, 2011.
- [8] Maximilian Kohlbrenner, Russell Hofmann, Sabbir Ahmed, and Youssef Kashef, “Pre-training cnns using convolutional autoencoders,” .
- [9] Matthew D Zeiler, Dilip Krishnan, Graham W Taylor, and Rob Fergus, “Deconvolutional networks,” in *Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on*. IEEE, 2010, pp. 2528–2535.
- [10] Wenzhe Shi, Jose Caballero, Lucas Theis, Ferenc Huszar, Andrew Aitken, Christian Ledig, and Zehan Wang, “Is the deconvolution layer the same as a convolutional layer?,” *arXiv preprint arXiv:1609.07009*, 2016.
- [11] Vincent Dumoulin and Francesco Visin, “A guide to convolution arithmetic for deep learning,” *arXiv preprint arXiv:1603.07285*, 2016.
- [12] J Zhao, M Mathieu, R Goroshin, and Y Lecun, “Stacked what-where auto-encoders. arxiv 2015,” *arXiv preprint arXiv:1506.02351*.
- [13] Karol J. Piczak, “Esc: Dataset for environmental sound classification,” in *Proceedings of the 23rd ACM International Conference on Multimedia*, New York, NY, USA, 2015, MM ’15, pp. 1015–1018, ACM.
- [14] Brian McFee, Matt McVicar, Oriol Nieto, Stefan Balke, Carl Thome, Dawen Liang, Eric Battenberg, Josh Moore, Rachel Bittner, Ryuichi Yamamoto, Dan Ellis, Fabian-Robert Stoter, Douglas Repetto, Simon Waloschek, CJ Carr, Seth Kranzler, Keunwoo Choi, Petr Viktorin, Joao Felipe Santos, Adrian Holovaty, Waldir Pimenta, and Hojin Lee, “librosa 0.5.0,” Feb. 2017.
- [15] “Pescador,” Retrieved January 2, 2018 from <https://github.com/pescadores/pescador>.
- [16] Hyeonwoo Noh, Seunghoon Hong, and Bohyung Han, “Learning deconvolution network for semantic segmentation,” in *Proceedings of the IEEE International Conference on Computer Vision*, 2015, pp. 1520–1528.
- [17] Yandre MG Costa, Luiz S Oliveira, and Carlos N Silla, “An evaluation of convolutional neural networks for music classification using spectrograms,” *Applied Soft Computing*, vol. 52, pp. 28–38, 2017.
- [18] Shawn Hershey, Sourish Chaudhuri, Daniel PW Ellis, Jort F Gemmeke, Aren Jansen, R Channing Moore,

Manoj Plakal, Devin Platt, Rif A Saurous, Bryan Seybold, et al., “Cnn architectures for large-scale audio classification,” in *Acoustics, Speech and Signal Processing (ICASSP), 2017 IEEE International Conference on*. IEEE, 2017, pp. 131–135.

- [19] Dumitru Erhan, Yoshua Bengio, Aaron Courville, Pierre-Antoine Manzagol, Pascal Vincent, and Samy Bengio, “Why does unsupervised pre-training help deep learning?,” *Journal of Machine Learning Research*, vol. 11, no. Feb, pp. 625–660, 2010.